

# Transient Heat Equation in Cylindrical Coordinates

Justin Miller  
APMA2822b

# Motivation:

- Background in power transmission
- Wire Overheating: Sag Clearance & Annealing
- Knowing wire temperature under different loading is valuable to optimal grid operation
- Emergency Loading Situations
  - Abrupt high-current power flow to fill demand until the conductor reaches Max. Rated Temp.



# Equations:

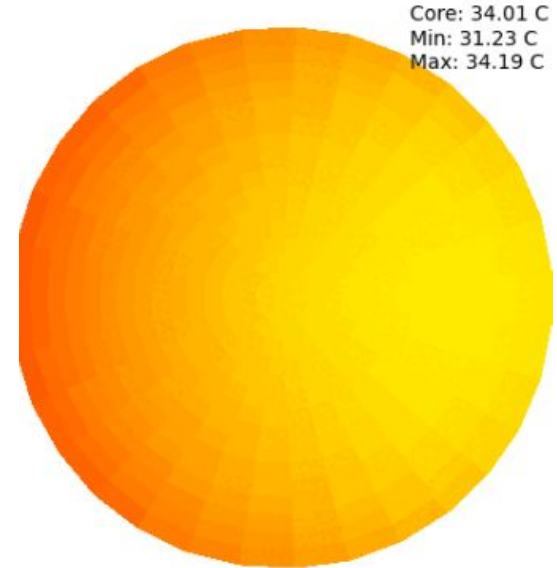
- Fourier-Biot Equation

$$\frac{\partial^2 T}{\partial x^2} + \frac{\partial^2 T}{\partial y^2} + \frac{\partial^2 T}{\partial z^2} + \frac{\dot{e}_{\text{gen}}}{k} = \frac{1}{\alpha} \frac{\partial T}{\partial t}$$

- Cylindrical Version

$$\frac{1}{r} \frac{\partial}{\partial r} \left( kr \frac{\partial T}{\partial r} \right) + \frac{1}{r^2} \frac{\partial T}{\partial \phi} \left( k \frac{\partial T}{\partial \phi} \right) + \dot{e}_{\text{gen}} = \rho c \frac{\partial T}{\partial t}$$

- Note: Using a 2D Cross Section of the above equation ( $dT/dz = 0$ )



# Finite Difference Scheme:

- Spatial Derivatives Appx. Using Central Difference
- Time Derivative Appx. Using Backward Difference

$$\frac{1}{r} \frac{\partial}{\partial r} \left( kr \frac{\partial T}{\partial r} \right) + \frac{1}{r^2} \frac{\partial T}{\partial \phi} \left( k \frac{\partial T}{\partial \phi} \right) + \dot{e}_{\text{gen}} = \rho c \frac{\partial T}{\partial t}$$

$O(\Delta x^2)$  centered difference approximations:

$$f'(x) : \quad \{f(x + \Delta x) - f(x - \Delta x)\} / (2\Delta x)$$

$$f''(x) : \quad \{f(x + \Delta x) - 2f(x) + f(x - \Delta x)\} / \Delta x^2$$

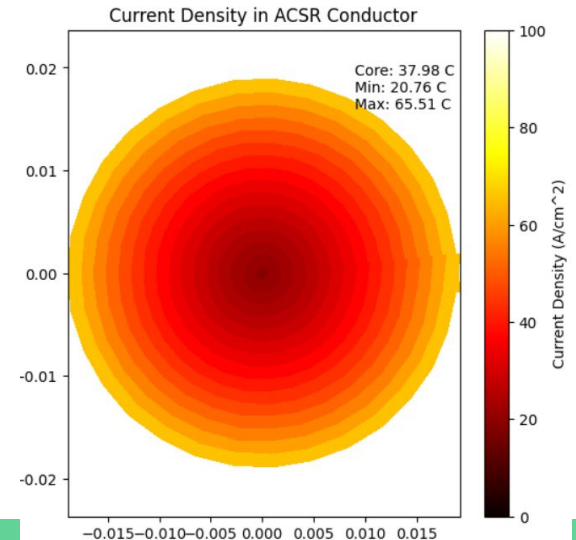
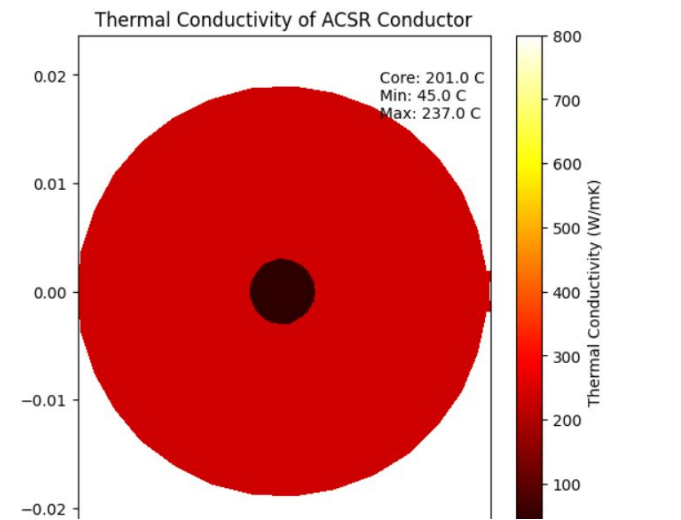
$O(\Delta x^2)$  backward difference approximations:

$$f'(x) : \quad \{3f(x) - 4f(x - \Delta x) + f(x - 2\Delta x)\} / (2\Delta x)$$

- Challenges:
  - Singularity Handling
  - Uneven Domain (much more time resolution required for this problem of interest)
- Benefits
  - Stability
  - Ease of calculation / parallelization

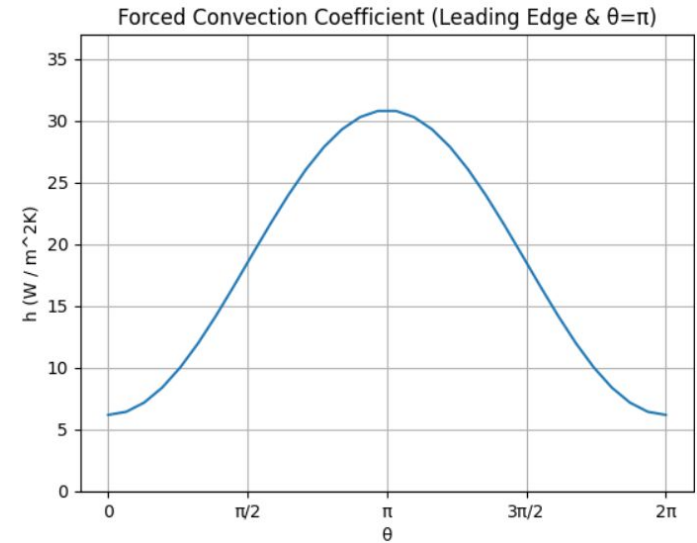
# Model Parameters

- “ACSR Linnet” Transmission Wire
- 1.84 cm radius; Aluminum w/ Steel Core
- K, alpha material properties
- Joule Heating based off of Current Load
  - Adjusted for skin effect




# Boundary Conditions

- Forced Convection:
  - 2 m/s perpendicular wind, constant direction
  - 25 °C Ambient Temperature°
  - Results in **asymmetric** convection pattern
    - Key reason for FD model



- Periodic Boundary Condition
  - Ensure wraparound in theta dir

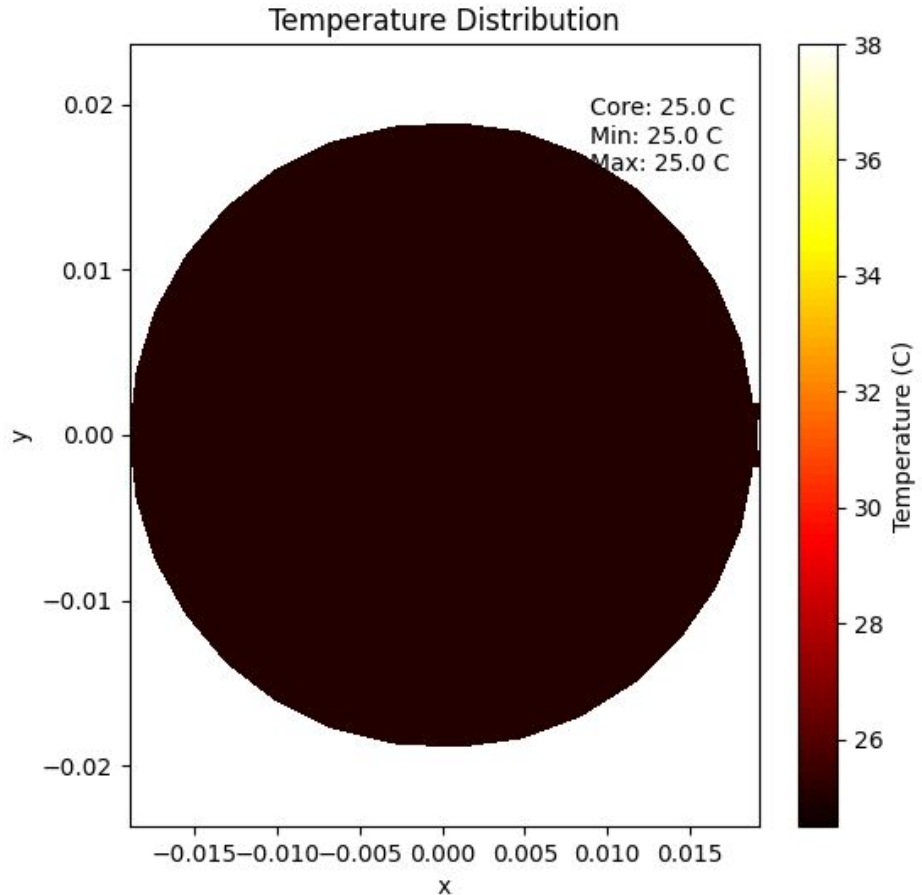


# Domain Parameters

- $N_r = 32$
- $N_{\theta} = 64$
- $N_t = 512$ 
  - $\Delta t = 10$  seconds
  - Majority of resolution is in  $t$  direction to simulate  $\sim$  an hour and a half of transient heating
- Initial Conditions:
- Power line at ambient temperature (25 C)
- Internal Heat Generation started at first timestep.

# Sequential Runs: Visualizing

- Transient heating of wire shown
  - Takes about an hour to reach steady state
- Uneven Convection Apparent
- Overall Success
  - Smooth heat transfer over singularity
  - Convection at surface as expected
  - Symmetric across horizontal plane
  - No Halo remnants in parallel code






## Misc. Notes: Singularity Handling & Convergence:

- Singularity at  $i=0$  ( $R=0$ )
- Effectively skip cell, take difference formulas from  $i+1$  on the current theta and the theta across the center
- Periodicity: treat last element of theta as before first element

- L1 Convergence;  
Tolerance:  $1e-6$
- 42,000 iterations required, consistent among all methods
- Stable, monotonically decreasing

# Parallelizing: Shared Overview

- Create `FiniteDifferenceKernel()` to run every iteration
- 3D Simulation: Splitting into 3D Threads / Blocks
- Allocating & initializing data to device using `Memcpy()`
- Execute kernel on each iteration.
  - Sync after run
  - Check convergence if desired
  - `Memcpy()` back to host when done



# Parallelizing: Shared Optimization

- Hipcc built in optimization -O2
- Thread Management
  - Recall domain size is N\_R:32, N\_theta:64, N\_t:512
  - By splitting our threads into (2, 4, 16) we can keep our grid sizes nice and even

# Parallelizing: Shared Performance

- Run on MI2104x Partition, Single GPU
- MI210: Stated 22.6 TF & up to a 1.6 TB/s BW
- 53.9% Theoretical Gf/s
- 54.0% Theoretical Gb/s

## Kernel Computation:

for each  $i * j * k$ : // **Iteration**

3 store;  
8 loads  
60 FLOP

**Per Iteration:**  $ijk \times 8 \times 11 = 88$   $ijk$  bytes,  
 $ijk \times 60$  FLOP = 60  $ijk$  FLOP,  
 $AI = 60 / 88$   
**Performance = 1091 Gf/s**

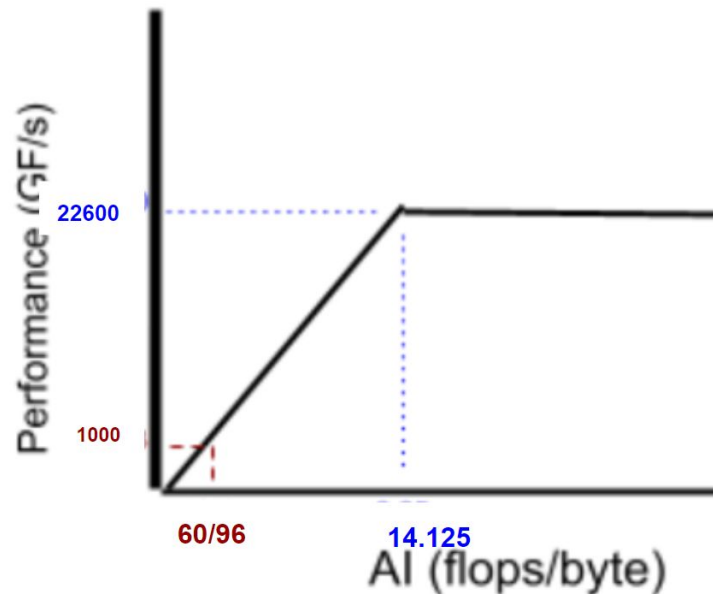
**Gf / s = 587.986**

**Gb / s = 862.380**

Tolerance	Runtime (s)	Iterations	Time / Iteration ( $\mu$ s)	Bytes / Iteration	Gb/s	FLOP / Iteration	Gf/s
1e-6	3.946	37,000	107	92,274,688	862.380	62,914,560	587.986

## Parallelizing: Shared Performance (cont.)

- Performance roofline plot provided
- % of theoretical efficiency is a bit worse than expected
- Practically acceptable



Tolerance	Runtime (s)	Iterations	Time / Iteration ( $\mu$ s)	Bytes / Iteration	Gb/s	FLOP / Iteration	Gf/s
1e-6	3.946	37,000	107	92,274,688	862.380	62,914,560	587.986

# Parallelizing: Distributed (offloading onto GPUs)

- Executing same kernel distributed on 4 MI210 GPUs (using 4 MPI Tasks)
- Using the rankID of each task, set a corresponding GPU device to the CPU
- Leveraging uneven domain to simplify and speed up halo exchange
  - **36 x 64 x 512 simulation: splitting this into 4 GPUs the lowest data transfer is in the t plane**
  - **Backwards difference in time plane simplifies it further, as buffer is only needed on one side**

# Parallelizing: Distributed Optimization

- Convergence Checks
  - Limiting convergence checks to every N iterations (ended up with 1000) to limit reductions
- Halo Exchange
  - Also limiting halo exchange to every N iterations as this is expensive. Slightly increases iterations needed for convergence but results in overall runtime improvements.
- Singularity
  - There are no extra considerations required to handle the singularity in the distributed model as we are splitting only on the time plane

# Overall Performance:

- Python Implementation, C++ Sequential Implementation, Shared, Distributed

	<b>Time / Iteration (<math>\mu</math>s)</b>	<b>Evaluated Runtime (s)</b>
Python Prototype	94,391	65 Min.
C++ Sequential	40,918	28 Min.
C++ GPU Shared	107	3.95 Sec.
C++ GPU Offloaded (4 GPUs)	89	3.30 Sec



# Issues

- Performance Issues: Particularly on shared memory
- Unsolved: Joule Heating Stability. When input current is too high the simulation diverges. Likely a bug in the current distribution.
- Time Prioritization

Questions?